MS&E 125: Intro to Applied Statistics

Linear regression

Professor Udell

Management Science and Engineering
Stanford

April 17, 2023

Outline

Linear models

Prediction

Fitting linear regression

Maximum likelihood

Multiple regression

Motivation: linear models

Linear models can be used for

- **prediction:** given a set of input variables, predict a value for the output variable
- **understanding:** how are the input variables related to the output variable, and to each other?
- **inference:** how much do the input variables affect the output variable?
- **counterfactuals:** what would happen if we changed the input variables?
- **control:** how can we change the input variables to achieve a desired output?

Outline

Linear models

Prediction

Fitting linear regression

Maximum likelihood

Multiple regression

Regression setup

we want to predict output given inputs

- input variables $x\mathbf{R}^p$
  - also called "predictors", "independent variables", "covariates"
  - a row of a data table
- output variable $y \in \mathbf{R}$
  - also called "outcome", "response", "dependent variable", "label", "target" ...

Regression setup

we want to predict output given inputs

- input variables $x\mathbf{R}^p$
  - also called "predictors", "independent variables", "covariates"
  - a row of a data table
- output variable $y \in \mathbf{R}$
  - also called "outcome", "response", "dependent variable", "label", "target" ...

example: to predict the cost of an insurance claim,

- y is the cost of an insurance claim.
- entries of x are the properties of the insured and his/her vehicle, e.g., credit score, age of the vehicle, ...

Demo: simple linear regression

https://colab.research.google.com/github/stanford-mse-125/demos/blob/main/regression.ipynb

Simple linear regression

simple linear regression: $p = 1$

- predict
$$\hat{y} = \beta_0 + \beta_1 x$$
- $\beta_0, \beta_1 \in \mathbf{R}$ are called **regression coefficients**
- $\hat{y}$ is called the **prediction** for input x

Predictions: example

In the fathers and sons dataset, we found

$$\hat{y} = 34 + 0.5x$$

where x is the height of the father in inches.

Predictions: example

In the fathers and sons dataset, we found

$$\hat{y} = 34 + 0.5x$$

where x is the height of the father in inches.

Q: What do the numbers 34 and .5 mean?

Predictions: example

In the fathers and sons dataset, we found

$$\hat{y} = 34 + 0.5x$$

where x is the height of the father in inches.

Q: What do the numbers 34 and .5 mean?
A: A father with height 0 inches has a son with height 34 inches. For each inch of height, the son is expected to be 0.5 inches taller.

Outline

Linear models

Prediction

Fitting linear regression

Maximum likelihood

Multiple regression

Residuals

look at **residual** r to understand how well the model fits the data

$$r = y - \hat{y} = y - \beta_0 - \beta_1 x_1$$

pick β so the residuals are small

Dataset

to find the best line, we need a dataset! suppose we have

- n data points $(x^{(1)}, y_1), \ldots, (x^{(n)}, y_1)$
  - also called **dataset**, **examples**, **observations**, **samples** or **measurements**
- each $x_i \in \mathbf{R}^p$ is a vector of p input variables
  - a

  row from the data table
- each $y_i \in \mathbf{R}$ is a scalar output variable

Linear regression: two perspectives

how to choose β?

- **optimization perspective:** find β to minimize the sum of squared errors

$$\text{minimize} \qquad \sum_{i=1}^{n} (y_i - \hat{y}_i)^2$$

- **statistical perspective:** find the line that maximizes the likelihood of the data

theorem: for appropriate assumptions, the two perspectives give the same answer
(coming in a few slides, or see All of Statistics ch. 14)

Least squares fitting

Least squares fitting

Q: Suppose $x_i = 0$ for every i. What is β_0?

Least squares fitting

Q: Suppose $x_i = 0$ for every i. What is β_0?
A: Set derivative to zero; solution is the average of the y_is.

Least squares fitting

Q: Suppose $x_i = 0$ for every i. What is β_0?
A: Set derivative to zero; solution is the average of the y_is.
Q: Given $x_i \in \mathbf{R}$, what is β_1?

Least squares fitting

Q: Suppose $x_i = 0$ for every i. What is β_0?
A: Set derivative to zero; solution is the average of the y_is.
Q: Given $x_i \in \mathbf{R}$, what is β_1?
A: Set derivative to zero; solution is the slope of the line of best fit.

Solve for β_0

$$\text{minimize} \qquad \sum i = 1^n (y_i - \beta_0 - \beta_1 x_i)^2$$

take derivative wrt β_0 and set to zero:

$$\begin{aligned}
\sum_{i=1}^{n} -2(y_i - \beta_0 - \beta_1 x_i) &= 0 \\
\sum_{i=1}^{n} y_i &= \beta_0 n - \beta_1 \sum_{i=1}^{n} x_i \\
\frac{1}{n}\sum_{i=1}^{n} y_i &= \beta_0 - \beta_1 \frac{1}{n}\sum_{i=1}^{n} x_i
\end{aligned}$$

$\Longrightarrow$ the model goes through the point of averages

Solve for β_1

$$\text{minimize} \qquad \sum i = 1^n (y_i - \beta_0 - \beta_1 x_i)^2$$

take derivative wrt β_1 and set to zero:

$$\begin{aligned} \sum_{i=1}^{n} -2(y_i - \beta_0 - \beta_1 x_i) x_i &= 0 \\ \sum_{i=1}^{n} x_i y_i &= \beta_0 \sum_{i=1}^{n} x_i + \beta_1 \sum_{i=1}^{n} x_i^2 \end{aligned}$$

$$\beta_1 = \frac{\sum_{i=1}^n x_i y_i - \beta_0 \sum_i = 1^n x_i}{\sum_{i=1}^n x_i^2}$$

interpretation:

- suppose x and y have been standardized so that $\sum_{i=1}^n x_i = \sum_{i=1}^n y_i = 0$ and $\frac{1}{n}\sum_{i=1}^n x_i^2 = \frac{1}{n}\sum_{i=1}^n y_i^2 = 1$.
- then $\beta_1 = \frac{1}{n}\sum_{i=1}^n x_i y_i$ is the **correlation** between x and y

Outline

Linear models

Prediction

Fitting linear regression

Maximum likelihood

Multiple regression

Linear regression model

probabilistic model for linear regression: suppose the xs are fixed, and ys are generated by

$$y_i = \beta_0 + \beta_1 x_i + \epsilon_i$$

where ϵ_i is a random variable with mean 0 and variance σ^2.

Linear regression model

probabilistic model for linear regression: suppose the xs are fixed, and ys are generated by

$$y_i = \beta_0 + \beta_1 x_i + \epsilon_i$$

where ϵ_i is a random variable with mean 0 and variance σ^2.

under the model, the likelihood of observing residual $r = y - \hat{y}$ is

$$\frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{r^2}{2\sigma^2}\right)$$

Demo: are errors iid normal?

https://colab.research.google.com/github/stanford-mse-125/demos/blob/main/regression.ipynb

Maximum likelihood

the **likelihood function** gives the probability of the data given the parameters:

$$\ell(\beta_0, \beta_1) = \prod_{i=1}^{n} \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y_i - \beta_0 - \beta_1 x_i)^2}{2\sigma^2}\right)$$

maximum likelihood chooses β_0 and β_1 to maximize the likelihood function:

$$\begin{aligned}
\hat{\beta}_0, \hat{\beta}_1 &= \operatorname*{argmax}_{\beta_0, \beta_1} \ell(\beta_0, \beta_1) \\
&= \operatorname*{argmax}_{\beta_0, \beta_1} \log \ell(\beta_0, \beta_1) \\
&= \operatorname*{argmax}_{\beta_0, \beta_1} \sum_{i=1}^{n} \log \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y_i - \beta_0 - \beta_1 x_i)^2}{2\sigma^2}\right) \\
&= \operatorname*{argmax}_{\beta_0, \beta_1} -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^{n} (y_i - \beta_0 - \beta_1 x_i)^2 \\
&= \operatorname{argmin} \frac{1}{\ } \sum^{n} (y_i - \beta_0 - \beta_1 x_i)^2
\end{aligned}$$

Probabilistic interpretation

$$y_i = \beta_0 + \beta_1 x_i + \epsilon_i, \qquad \epsilon_i \sim \mathcal{N}(0, \sigma^2)$$

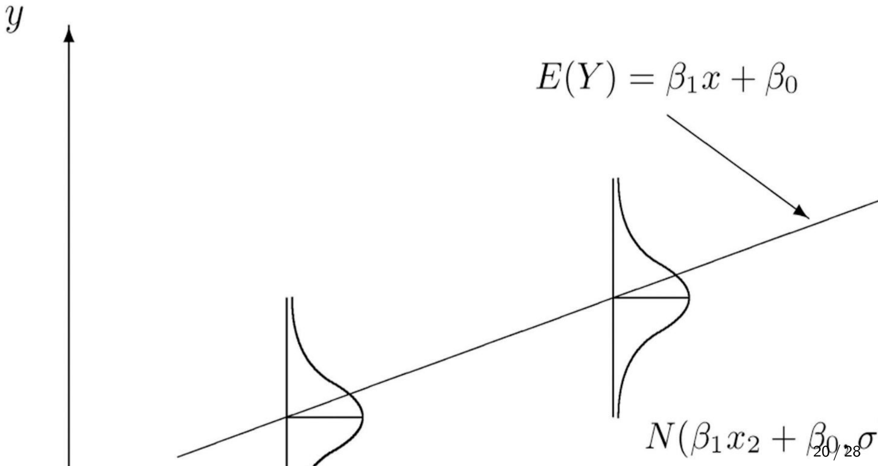

Estimation puts a hat on it

statisticians use hats to denote estimates:

- $\hat{\beta}_0$ is the estimate of β_0
- $\hat{y}$ is the estimate of y

these estimates are **random quantities** that depend on the data

`https://colab.research.google.com/github/stanford-mse-125/demos/blob/main/regression-uncertainty.ipynb`

Properties of the estimator

putting it together, we have found:

$$\hat{\beta}_1 = \rho(x, y)\hat{\sigma}_y/\hat{\sigma}_x, \qquad \hat{\beta}_0 = \bar{y} - \hat{\beta}_1\bar{x}$$

where

- $\rho(x, y)$ is the correlation between x and y
- $\hat{\sigma}_x$ and $\hat{\sigma}_y$ are the sample standard deviations of x and y
- $\bar{x}$ and $\bar{y}$ are the sample means of x and y

under the normal model

$$y_i = \beta_0 + \beta_1 x_i + \epsilon_i, \qquad \epsilon_i \sim \mathcal{N}(0, \sigma^2)$$

these estimates are unbiased:

$$\mathbb{E}[\hat{\beta}_1] = \beta_1, \qquad \mathbb{E}[\hat{\beta}_0] = \beta_0$$

see All of Statistics ch 14 for formulas for the variance of the estimates

Outline

Linear models

Prediction

Fitting linear regression

Maximum likelihood

Multiple regression

Matrix notation

$$y_i = \beta_0 + \beta_1 x_i + \epsilon_i, \qquad \epsilon_i \sim \mathcal{N}(0, \sigma^2)$$

rewrite using linear algebra:

- form **response vector** $y \in \mathbf{R}^n$: each outcome y_i is an entry of y
  - a

  lso called **target vector**
- form **design matrix** $X \in \mathbf{R}^{n \times p}$: each example $x^{(i)}$ is a row of X
  - also called **feature matrix**
  - if the model includes a constant term, the 0th column of $X \in \mathbf{R}^{n \times p+1}$ is all ones

$$\begin{bmatrix} y_1 \\ \vdots \\ \vdots \\ y_n \end{bmatrix} = \begin{bmatrix} X_{11} & \cdots & X_{1p} \\ \vdots & \vdots & \vdots \\ \vdots & \vdots & \vdots \\ X_{n1} & \cdots & X_{np} \end{bmatrix} \begin{bmatrix} \beta_1 \\ \vdots \\ \beta_p \end{bmatrix} + \begin{bmatrix} \epsilon_1 \\ \vdots \\ \vdots \\ \epsilon_n \end{bmatrix}$$

Least squares in matrix notation

rewrite error:

$$\sum_{i=1}^{n}(y_i - \hat{y}_i)^2 = \|y - X\beta\|^2$$

interpretation:

- $X\beta$ is a linear combination of the columns of X
- we seek the linear combination that best matches y

Linear regression: model

we can rewrite the model as

$$y_i = \beta_0 + \beta_1 X_{i,1} + \beta_2 X_{i,2} + \cdots + \beta_p X_{i,p} + \varepsilon_i$$

- notice that $\beta_0, \beta_1, \ldots, \beta_p$ do not depend on i
- the columns of the data table are $Y_i, X_{i,1}, \ldots, X_{i,p}$

i	y_i	$X_{i,1}$	$X_{i,2}$	...	$X_{i,p}$
1	2.3	1.1	6.2	...	5.9
2	12.7	2.4	5.4	...	9.6
3	6.3	0.9	6.9	...	1.5

Example: electricity usage

- We are managing a large complex of apartments in the Northeast.
- We pay for the electricity used by our residents.
- We would like to predict electricity usage so that we can estimate how much money should be set aside.

Demo: multiple linear regression

```
https://colab.research.google.com/github/
stanford-mse-125/demos/blob/main/electricity.ipynb
```